



Brian Bagdasarian

Software Architect

Los Angeles, CA bbagdasarian.scs@gmail.com

[bbagdasarian](#)

[ahappyypie](#)

Summary

I'm a systems designer, platform builder, and process improver. I've led transformation and construction of large scale data systems, internal developer platforms, and order-of-magnitude process efficiency increases.

I am seeking a strategic leadership role where I can make a long-term impact with an organization.

Education

New York University

September 2013 - December 2016

Gallatin School of Individualized Study, BA in Data Driven Decision Making
Courant Institute of Mathematics, Joint Minor in Computer Science & Mathematics

Experience

Amberdata

Senior Data Platform Engineer

July 2024 - February 2025

Remote

Decreased alerts and incidents upwards of 80% and increased performance by 250% by analyzing continuous profiling of key backend services.

Provided architectural guidance to teams working on moving dozens of pipelines and petabytes of data from batch to streaming processes.

Improved developer platform to increase product velocity from weekly to multiple times per day without violating SLAs.

Key technologies: Apache Spark, Delta Lake, Databricks, Redpanda, Apache Pinot, Kubernetes, ArgoCD, Pyroscope, AWS, Terraform, Bazel

Cosm

Lead Platform Engineer

September 2022 - June 2024

Playa Vista, CA | Hybrid

Founding member of Cosm's Platform Team, dedicated to improving developer experience, reliability and performance of Cosm's core cloud and on premise infrastructure.

Designed and implemented a custom infrastructure as code framework with Pulumi, which allowed for granular control of cloud and on premise resources.

Built Cosm's network and Kubernetes infrastructure, along with other platform engineers. Responsible for integration of EKS and supporting tools, notably Cilium and Flux.

Extended platform Kubernetes clusters to support data infrastructure, including integration between Redpanda, Strimzi and Debezium for change data capture.

Led cross functional integration efforts between several departments, including finance and venue operations for complete data integration with Databricks.

Led Cosm's CPRA & GDPR compliance efforts through integration with Transcend, Cosm's website, mobile app and backend.

Managed the data team, driving product requirements and delivery along with stakeholders in other departments.

Participated in Cosm's internship program as a manager in multiple cohorts, mentoring and training engineers, including converting an intern to a full time engineer at the conclusion of the program.

Key technologies: AWS, Kubernetes, Cilium, Flux, Pulumi, Redpanda, Strimzi, Debezium, Apache Spark, Delta Lake, Databricks, S3, Transcend

Cosm

Senior Data Engineer

September 2021 - September 2022

Playa Vista, CA | Hybrid

Designed a Data Lakehouse in collaboration with stakeholders using Medallion Architecture to centralize, govern and deliver Cosm's disparate venue and business data. Critical data sources were delivered via real time streaming and transformed using incremental processing.

Key technologies: Apache Spark, Delta Lake, Databricks, Redpanda, Debezium, AWS, Pulumi

Nativo

Big Data Engineer

January 2020 - September 2021

El Segundo, CA | Hybrid

Designed next generation streaming data architecture to meet the analytics product SLA by reducing metric arrival latency from hours to minutes, while increasing data volume 100% and reducing pipeline cost 20% using a transactional data lake.

Optimized data warehousing infrastructure, increasing typical query performance 20% and reducing cost 40%.

Key technologies: Apache Spark, Delta Lake, Databricks, Snowflake, Kubernetes, AWS, Terraform

Regent

Data Engineer

June 2018 - December 2019

Los Angeles, CA | On Site

Designed and built a shared services data platform, greatly increasing the speed at which new acquisitions could be onboarded, a critical metric of restoring companies to profitability.

Planned and executed retail software platform migration across 800+ locations, including point of sale and employee authentication, resulting in an 80% decrease in day to day management tasks compared to the previous solution.

Key technologies: Apache Spark, MySQL, Snowflake, Kubernetes, AWS, DigitalOcean, Terraform

FIJI Water

Data Engineer

April 2017 - May 2018

Los Angeles, CA | On Site

Architected and built FIJI Water's first automated data pipeline and business intelligence platform, saving 600+ hours each week on ad hoc depletions reporting and competitive set analysis.

Led trainings for the nationwide sales team, focused on self service business intelligence, allowing representatives on demand access to verified reports and analysis across desktop and mobile environments.

Key technologies: Apache Spark, Microsoft SQL Server, Microsoft Azure, Microsoft PowerBI