

Brian R. Bagdasarian

Web: ahappypie.github.io

Email: bagdasarian.scs@gmail.com

Technical Experience

Currently In Production

Platform & Infrastructure

Infrastructure: Pulumi, AWS

Orchestration: Kubernetes

Delivery: Helm, Kustomize, FluxCD, Garden,

GitHub Actions

Observability: OpenTelemetry, Grafana, Mimir,

Loki, Tempo

Additional Services: Temporal, Tailscale

Data

Database: MongoDB Atlas

Compute: Databricks

Streaming: Redpanda

Storage: S3 + Delta Lake

Languages

TypeScript: Pulumi, Frontend

Rust: Backend

RPC: gRPC

Scala: Apache Spark

SQL: Spark, Databricks SQL

Priority Research

nix, Cilium, Istio, Apache Arrow

Past Experience

Java (8), PostgreSQL, MySQL, SQL Server, Apache Kafka, NATS, Snowflake, Microsoft Azure, DigitalOcean, Google Cloud Platform

Open Source

[ahappypie/sol](#)

Kubernetes reference architecture using CNCF tools

[ahappypie/marschat](#)

API for sending messages based on Earth-Mars light delay

[ahappypie/unica](#)

Generate unique identifiers at high scale. Based on Twitter's Snowflake project

[ahappypie/spotter](#)

Experiments with Spark, Kafka, Avro, Schema Registry, Hudi and others

Education

New York University

Gallatin School of Individualized Study

Courant Institute of Mathematics

Graduated Fall 2016

BA, Politics & Data Science

Joint Minor, Computer Science & Mathematics

Experience

Cosm – Lead Platform Engineer

September 2022 – Present

Helping design, build and maintain a multi-language monorepo for fully remote development
Designed infrastructure-as-code framework for granular control over cloud resources
Proving immutable infrastructure concepts for legacy on-premise systems

Cosm – Senior Data Engineer

September 2021 – September 2022

Designed data systems for powering next generation immersive experiences

Nativo Inc. – Big Data Engineer

January 2020 – September 2021

Responsible for new streaming data architecture for 5+ TB of new data per day using transactional data lake
Assisted in managing data warehouse of 150+ TB

Regent, L.P. – Data Engineer

July 2018 – December 2019

Restored companies to profitability with cloud native technologies
Engineered data platform for shared services
Responsible for maintenance of data warehouse
Planned and executed retail platform migration across hundreds of physical locations

FIJI Water & The Wonderful Company - Data Engineer

April 2017 – May 2018 (full-time)
Summer 2015 & 2016 (intern)

Architected and developed ETL pipeline and data warehouse
Led self-service analytics, including training courses
Consulted on reporting and contract management systems

BlueLine Grid - Software Engineer

February 2014 - September 2014

Supported iOS app deployment, developed backend systems with Akka